On File

Round 1 proposal, UGREEN Hackathon 2026 on AI NAS. Track: Open Innovation.

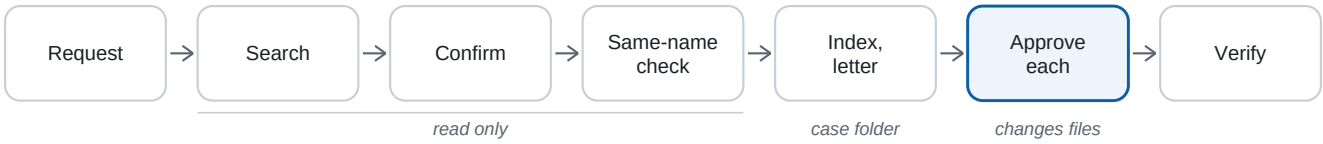

1. What does your Skill do?

On File finds files in chosen NAS folders mentioning one patient or customer, quotes the passage, drafts the legal reply, and changes files only after per-file approval.

2. Who is your target user, and what is their problem?

Small practices with forms, scans, photos and voicemails on a NAS. When someone requests everything held on them, US law (HIPAA) allows 30 days and EU law (GDPR) one month, no small-business exemption. In 2026 Shen Smiles, a ten-person dental practice, paid 140,000 USD after one patient's records request went unanswered.

3. What is your solution?

It searches file names, text, scanned pages, photo text and captions, document properties, voicemails, subtitles and camera face names. It separates same-name people and extracts the person's entries from shared files. The Skill itself sends nothing off the NAS.

4. What is the main user flow?

Input: request type, name, date received, folders, case folder, report-only or act mode. It shows days left and three lists (included, needs your decision, unreadable), then saves an index and letter. In act mode each copy, extract or deletion waits for its own approval.

5. How is your idea different from existing tools or common AI solutions?

PII Crawler (497 USD, local) already finds one person's data, scans included; OneTrust and Transcend automate requests in business systems; chat assistants need uploads. The difference: no model alone includes or excludes a file; identifiers and quotes checked against the file decide, so a same-name mother stays out. It then finishes: extracts, retained records, letter, per-file approval.

6. Which UGREEN AI NAS features does your Skill use, and at which step?

file-agent: find, cat, stat, sha256sum, scan-dedupe, share-mkdir, mkdir, cpmv, del, echo, restore, tag-add, tag-add-file (steps 0, 1, 3, 6, 7, 8). knowledge-agent: collect-evidence, audio-transcription, save-text (2, 3, 6, 7). document-agent: qa, metadata, extract, writing, translate (3 to 7). photo-agent: search-picture, vision-ocr, vision-vlm, meta (3 to 5). media-agent: search (3). security-monitor-agent: faces (3).

7. What is your technical approach?

It directs six NAS expert agents. Only step 7 changes the person's files, one approval each. A name counts only beside the person's identifiers; name-only files get for and against checks; unclear ones go to the user. Erasure cases keep a record, not the erased data.

8. What have you built so far? What is not built yet? What still needs testing on a NAS?

Built: SKILL.md and eight references, checker-clean; three synthetic practices with answer keys, one written blind by an agent that never saw the rules. Laptop stand-in (small models, name only, real audio): 40 of 40 and 43 of 43 requester files found, none wrongly included, 0 of 48 look-alikes. Blind: 18 of 30 found at first, 29 of 30 after fixes, still none wrong. Not built: nothing planned. Needs a NAS: search recall, OCR, approvals, restore.

9. How will you test your Skill if you become a Finalist?

All three corpora. Core 1: every requester file found, shared files planned as extracts, letter written. Core 2, erasure: hash files, approve each deletion; only planned deletions gone, all else unchanged. Edge: the requester's mother, also Maria Lopez, same address; it asks once, then includes none of hers.

10. What are the known limits, risks, or outside dependencies of your Skill?

Nicknames need supplied identifiers. Poor scans can hide a name. Archives and protected files are flagged; mailboxes and practice software are out of reach. Camera identities and shared files are edited by hand. Models are a NAS setting. Not legal advice. No outside service.

Not counted toward the 500 words

Team: Ujwal Vanjare, individual entrant.

Third-party components: none inside the Skill package (Markdown only). Testing tools outside the package: Python 3.12 (PSF licence), python-docx (MIT), openpyxl (MIT), reportlab (BSD), Pillow (MIT-CMU), pypdf (BSD), requests (Apache 2.0), Ollama (MIT), Qwen3 4B and Qwen3-VL 2B instruct (Apache 2.0), Windows offline speech (System.Speech). Sources: official PyPI, GitHub and Ollama distributions.

AI tools used: Claude Code (Anthropic) for rules research, collision scans, drafting SKILL.md, the references and these answers, the corpus generators, the scorer, the laptop harness and reviews against the judging criteria; a separate Claude agent, kept from the rules, wrote the blind test corpus. The idea, request types and naming are the team's own; the team directed the work, reviewed the drafts and made the final choices.

Test corpora and scorer (optional): https://github.com/usv240/on-file

Word count of the ten answers: 483. Laptop numbers in item 8 are stand-in measurements and will be replaced by NAS results in Round 2.